

Portfolio Backtest Report: Black-Litterman with CNN-BiLSTM Views

1. Introduction

This report outlines a comprehensive portfolio construction and backtesting framework using the Black-Litterman model enhanced with views generated from a CNN-BiLSTM architecture. The goal is to outperform the Nifty 50 benchmark using systematic investor views and robust optimization.

2. Universe & Dataset

- **Universe:** Top 50 stocks in the Nifty 50 Index.
- **Data Source:** Yahoo Finance (via yfinance).
- **Frequency:** Daily adjusted prices.
- **Duration:** 5 years.

I used price, volume, and various technical indicators as features:

- Indicators: RSI, MACD, EMA, MA, Bollinger Bands, etc.
 - Derived Features: Returns, Volatility, Price & Volume Change, Momentum.
-

3. Feature Engineering

Each stock's dataframe is enriched with:

- **Returns:** Daily % change in closing price.
- **Volatility:** Rolling standard deviation.
- **Momentum:** Price momentum over 10 days.
- **Bollinger Band position:** Price location within bands.

Missing and infinite values were forward-filled and backward-filled, and residual NaNs were set to 0.

4. Model: CNN-BiLSTM Architecture

Purpose:

To generate 5-day forward expected returns and uncertainty estimates.

Input:

- A rolling window of 30 days (sequence length) of technical features.

Layers:

- **Input:** Shape = (30, 19 features)
- **Noise + Conv1D + MaxPooling + BatchNorm** (feature extraction)
- **Bidirectional LSTM** (temporal dependencies)
- **Dropout layers with MC Dropout enabled** (for uncertainty)
- **Dense output layer**

Loss Function:

- Mean Squared Error (MSE)

Optimization:

- Adam optimizer with early stopping and learning rate reduction on plateau.

MC Dropout:

- Enabled at training and inference to generate predictive distributions.
- Predictions sampled 50 times; uncertainty = standard deviation.

Seed Setup for Reproducibility:

```
SEED = 42
os.environ['PYTHONHASHSEED'] = str(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)
tf.keras.utils.set_random_seed(SEED)
tf.config.experimental.enable_op_determinism()
```

5. Black-Litterman Model

Inputs:

- **Views:** Expected return for each stock from the CNN-BiLSTM model.
- **Uncertainty:** Standard deviation of MC Dropout predictions.
- **Market Weights:** Derived from historical market caps (equal-weighted fallback).

Formula:

$$\text{Posterior Mean} = ((\tau\Sigma)^{-1} + P^T\Omega^{-1}P)^{-1} * ((\tau\Sigma)^{-1} * \Pi + P^T\Omega^{-1}Q)$$

Where:

- τ = scalar scaling parameter for uncertainty in prior covariance (set to 0.025).

- Π = implied returns from market equilibrium.
 - Q = CNN-BiLSTM generated views.
 - Ω = diagonal matrix of view variances.
-

6. Risk Aversion (λ) Calculation

Instead of a fixed λ , I dynamically compute it as:

$$\lambda = (E[R] - r_f) / \sigma^2$$

Where:

- $E[R]$: Market portfolio annualized return (Nifty 50).
- r_f : Risk-free rate (6%).
- σ^2 : Annualized variance of market returns.

This allows our optimizer to adjust to current market conditions.

7. Backtesting Framework

Type 1: Full Period Backtest

- Entire 5-year period.
 - One-shot training of models.
 - One optimization step.
- (Used only in debug phase)

Type 2: Out-of-Sample Biweekly Rebalancing

- First 60% for training.
- Remaining 40% for testing.
- Rebalance every 10 trading days.
- At each rebalance:
 - Slice latest available data.
 - Generate views.
 - Update market caps.
 - Recompute optimal weights.
 - Track portfolio returns.

Metrics Calculated:


- **Total Return**
- **Annualized Return**

- **Volatility**
- **Sharpe Ratio** (vs 6% risk-free rate)
- **Max Drawdown**
- **Cumulative Return Curve**

8. Results Summary

Type 2: Bi-Weekly Rebalanced Portfolio

Metric	Portfolio	Nifty 50	Excess
Annualized Return	22.49%	15.11%	7.38%
Volatility	13.76%	13.23%	—
Sharpe Ratio	1.198	0.689	—
Max Drawdown	-15.91%	-15.77%	—



This demonstrates a clear outperformance of the model-based portfolio over the benchmark, especially in risk-adjusted terms (Sharpe ratio).

Top Holdings:

Stock	Weight (%)
TITAN.NS	9.69%
RELIANCE.NS	7.58%
BHARTIARTL.NS	4.99%
M&M.NS	4.76%
ICICIBANK.NS	4.40%
INFY.NS	4.39%
ASIANPAINT.NS	3.88%
DIVISLAB.NS	3.67%
ADANIPORTS.NS	3.51%
CIPLA.NS	3.42%

9. Stress Testing

To assess robustness, I applied shocks to investor views and measured the impact on performance metrics over a 1-year test window.

Shocks Applied:

- **Bullish Shock:** +20% to all views
- **Bearish Shock:** -20% to all views
- **Single Asset Shock:** -20% to RELIANCE only

Returns in stress testing were computed **without rebalancing**, i.e., single-shot portfolio allocation and performance tracking through the period.

Results:

Annualized Return:

- Baseline: 22.13%
- Bullish Shock: 22.34%

- Bearish Shock: 21.90%
- Single Asset Shock (RELIANCE): 22.11%

Sharpe Ratio:

- Baseline: 0.993
- Bullish Shock: 1.004
- Bearish Shock: 0.979
- Single Asset Shock (RELIANCE): 0.991

The model demonstrates strong resilience across scenarios, with minor sensitivity to bearish shocks.

10. Sensitivity Analysis (τ Parameter)

The sensitivity analysis was initially designed to assess how different values of τ (the scaling parameter for uncertainty in the Black-Litterman model) influence portfolio behavior. However, the results showed highly inconsistent and fluctuating trends, indicating that model stability was not meaningfully improved by tuning τ .

Final Decision:

I selected $\tau = 0.025$, as it aligns with prior literature and ensures a balanced trade-off between incorporating the market equilibrium and respecting the investor views.

The sensitivity analysis section has been removed due to the erratic nature of results and lack of consistent improvement.

11. Limitations & Future Work

- Currently, τ is fixed based on empirical reasoning.
 - Sector-wise diversification not explicitly controlled.
 - Stress testing only modifies views; price shocks not incorporated.
-

12. Next Steps

- Investigate confidence-weighted view aggregation.
 - Incorporate fundamental data into view generation.
 - Expand stress testing to include market-wide and sector-specific price perturbations.
-

Appendix

A. Model Summary

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 30, 19)	0
gaussian_noise (GaussianNoise)	(None, 30, 19)	0
gaussian_noise (GaussianNoise)	(None, 30, 19)	0
conv1d (Conv1D)	(None, 30, 64)	3,712
spatial_dropout1d (SpatialDropout1D)	(None, 30, 64)	0
batch_normalizat... (BatchNormalizat...)	(None, 30, 64)	256
max_pooling1d (MaxPooling1D)	(None, 15, 64)	0

conv1d_1 (Conv1D)	(None, 15, 32)	6,176
spatial_dropout1...	(None, 15, 32)	0
batch_normalizat...	(None, 15, 32)	128
bidirectional	(None, 15, 100)	33,200
dropout (Dropout)	(None, 15, 100)	0
bidirectional_1	(None, 50)	25,200
dense (Dense)	(None, 50)	2,550
dropout_1	(None, 50)	0
dense_1 (Dense)	(None, 25)	1,275
dropout_2	(None, 25)	0

(Dropout)		
dense_2 (Dense)	(None, 1)	26

Total params: 72,529 (283.33 KB)

Trainable params: 72,331 (282.54 KB)

Non-trainable params: 192 (768.00 B)

Optimizer params: 6 (36.00 B)

B. Package Versions

Specified in “requirements.txt”

Note: All experiments were run with seeds and determinism enabled for reproducibility.